

Notes to Claude:

1. There are some hyperlinks included, please format it accordingly.

2. Title section in orange for highlight, same for the bold emphasis (you did it well).

3. There are some notes for you (since it's a txt/pdf file the formatting might be absent), one complete, please remove those notes.

4. when you see long quotes, please format it in the appropriate style.

Weekend with Claude: From Research to Reality

Closing My Own Capabilities Gap (section title, orange colour)

I learned Claude Code and Cursor on Saturday morning and shipped my complete website by Sunday afternoon. In between, I went for coffee and hit the gym—no panic, no rush, just inspired action fueled by quiet confidence in the tools I was learning to use. I decided to take the rate limit message "hit your limit · resets 4pm (Europe/Warsaw)" to be less a frustration and more a gentle forcing function for sustainable work. This weekend wasn't just about building a website; it was about experiencing firsthand the capabilities gap I've been seeing and researching theoretically—the distance between what AI tools can do and what users actually leverage.

[note: blog1_header_image1 goes here]

The Shift That Changed Everything (section title, orange colour)

It started with a conversation about applying to a Research Scientist position at Anthropic. Because, well, we use Claude for everything. My first instinct was excitement, but Claude's initial response was sobering:

"For a Research Scientist role at Anthropic, you're facing a significant gap. This position typically requires published research in ML/AI, demonstrated technical implementation skills, and a track record of original research contributions. Without formal technical projects or programming proficiency, this will be a very difficult application."

Fair enough. But then I asked a simple question: "What is the application deadline?"

What happened next was remarkable. Claude reconsidered—not because I defended myself or explained why I was qualified, but because it took a second look at the role itself:

"Looking at this specific role more carefully, it might actually be more compatible with your background than I initially thought... This completely changes the picture. You're actually a strong candidate for this specific role—you just don't realize it yet."

This pivot—this willingness to reevaluate initial assumptions—was a profound "aha moment" for me. It revealed something about Claude's architecture, the Constitutional AI framework underlying it, and the thoughtful approach of the people building it. The system didn't just follow a script but has reflected, reconsidered, and adapted. This change gave me confidence not just about the application, but about what I was about to attempt: learning new tools and building something real in a single weekend.

(And yes, my partner's encouragement helped too!)

Parallel Workflows: The New Atomic Habits (section title, orange colour)

Six years ago, I completed a full-stack development course. A couple years ago, I dabbled in creative coding. But it would have been a stretch to call myself a developer—even an amateur one. This weekend changed that.

What made the difference wasn't just the tools (Claude Code and Cursor), but the mental model shift they enabled. I could work on my website while simultaneously preparing application materials, perfecting my CV, and learning about the position. While Claude Code was building components or debugging, I was researching Anthropic's papers, refining my cover letter, or simply taking necessary breaks.

This is what I've been calling the "capabilities gap" in my research: the widening distance between what AI systems can do and how users actually leverage them. It's not just about knowing features exist—it requires a fundamental mindset shift. As [Nathan Lambert writes](#), we're still "getting good at agents," learning new ways of doing things. For me, that meant developing new atomic habits: "while I do this, the agent builds that" requires planning skills, project management, and trusting the system enough to work in parallel rather than sequentially.

Practical Insights: What I Learned

On tool selection: I quickly learned to match tasks to capabilities. Claude excels at reasoning and composition—complex problem-solving, architectural decisions, nuanced writing. For simpler, more mechanical tasks, I'd use ChatGPT or other tools (of course, I didn't tell GPT I'm applying to Anthropic to prevent any sabotage). This isn't about one being "better"—it's about understanding what each tool does best and orchestrating accordingly.

On Claude Code specifically: One surprising discovery—when I tried copying longer text passages to fill website sections, Claude Code would sometimes didn't get the complete version. The more effective strategy? Taking screenshots of the text and pasting those instead. A small workaround, but it reveals something about how these tools interpret different input formats.

On rate limits as design: The Pro plan's message limits became a feature, not a bug. They enforced breaks between working sessions, preventing burnout and giving me time to reflect on what I was building. Constraints can be generative.

From Theory to Practice: Living My Research (section title, orange colour)

The irony isn't lost on me: I study human-AI co-creativity, interaction patterns, and how users develop mental models for AI systems—and this weekend was among those moments when I experienced all of it firsthand. I became my own case study. In my master's thesis (note: doi [10.13140/RG.2.2.31725.86248](https://doi.org/10.13140/RG.2.2.31725.86248) for hyperlink), I analyzed co-creative practices with AI through HCI and art therapy frameworks, examining cognitive states activated during collaborative creation. I found that users prefer maintaining strong control while expecting substantial support from AI systems. Well, I believe I still managed to form a partnership while leading Claude toward my vision: I made all architectural decisions, but Claude handled implementation details, debugging, and optimization. The mixed-initiative balance worked because the tools respected my agency while augmenting my capabilities.

This experience also connects to emerging research directions I'm exploring. Dario Amodei recently identified "AI building AI" as a key research area for 2026. I just lived it—using Claude to build digital artifacts (my website), which itself required Claude to generate code, design systems, and solve technical problems. The human role shifted from implementation to orchestration, from coding to guiding, from building to designing.

What's Next: Creative Coding and Beyond (section title, orange colour)

I'm already thinking about what else becomes possible. Creative coding, for instance—I got super curious to test Claude Code for live VJ performances, producing visuals from code on-the-go. This wouldn't just be entertaining; it would be an interesting research project in itself. Using the interface (whether Claude Code, Cursor, or something else) as a mixed-initiative creative tool during live performance would reveal fascinating insights about real-time human-AI collaboration under time pressure.

More broadly, this weekend taught me that the capabilities gap isn't just an abstract research concept—it's something we all navigate daily. The tools exist. The challenge is developing the mental models, workflows, and "AI manners" (as I've been calling them, I am still guided by *the etiquette* when interaction with AI) to effectively collaborate with increasingly capable systems. It requires rethinking not just how we work, but how we think about work itself.

Many worry we're doomed to lose jobs as AI capabilities expand. I think we're facing something more interesting: a radical rethinking of how society organizes work, creativity, and knowledge production. If we take sustainable approaches—like I see Anthropic doing with frameworks like Constitutional AI among others, thoughtful rate limits, and emphasis on human-AI collaboration rather than replacement—we'll be fine. Actually, it could be better than fine, as we'll unlock possibilities and the world we haven't imagined yet.

The Meta-Reflection (section title, orange colour)

There's something beautifully recursive about using Claude to prepare an application to work on Claude. About using AI-assisted development to demonstrate technical capability. About experiencing the capabilities gap while researching it. About becoming both researcher and subject.

This weekend, I didn't just build a website. I closed my own capabilities gap. I shifted from studying human-AI interaction to embodying it. I moved from theoretical frameworks to practical implementation. And I gained confidence not just in the tools, but in my ability to learn, adapt, and ship—even under weekend time constraints.

The website is live. The application is submitted. And I'm already thinking about what to build next.

[A thin line in accent orange to separate the text, then in italic:]

This post is part of my application to Anthropic's Societal Impacts team. You can explore my research on human-AI interaction, AI governance, and cross-cultural technology studies at zoyaslavina.com.